

Topic: GPU+PyTorch+Triton Starting Tutorial

0. Hardware Information

Hardware: NVIDIA A100

Driver: cuda_12.3.r12.3/compiler.33281558_0

表：GPU硬件的信息汇总

架构	GPU型号	计算能力	-arch参数	备注	CUDA版本要求
Volta	V100	7.0	-arch=sm_70	数据中心卡	9.0+
Volta	Titan V	7.0	-arch=sm_70	消费级卡	9.0+
Ampere	A100	8.0	-arch=sm_80	数据中心卡	11.0+
Ampere	A30	8.0	-arch=sm_80	数据中心卡	11.0+
Ampere	RTX 30系列	8.6	-arch=sm_86	消费级卡	11.0+
Hopper	H100	9.0	-arch=sm_90	数据中心卡	11.8+
Blackwell	B200	9.0+	-arch=sm_90	目前使用Hopper参数	12.0+

1. Building Environment

```
wget https://repo.anaconda.com/miniconda/Miniconda3-latest-Linux-  
x86_64.sh  
bash ~/Miniconda3-latest-Linux-x86_64.sh
```

2. PyTorch + Triton + vLLM

```
pip3 install torch torchvision --index-url  
https://download.pytorch.org/whl/cu121
```

```
pip install triton==3.4.0
```

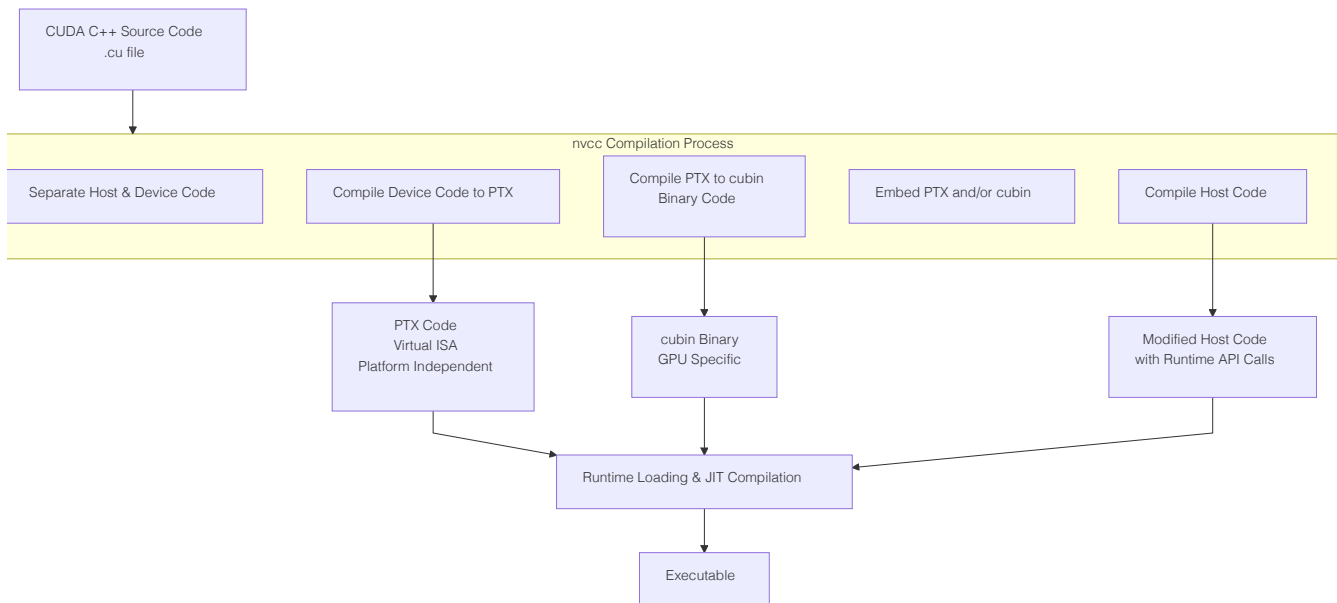
```
pip install vllm==0.9.2
```

Verify the installations:

```
# 验证脚本 verify_installation.py  
import torch  
import triton  
import vllm  
  
print(f"PyTorch 版本: {torch.__version__}")  
print(f"PyTorch CUDA 可用: {torch.cuda.is_available()}")  
print(f"PyTorch CUDA 版本: {torch.version.cuda}")  
print(f"Triton 版本: {triton.__version__}")  
print(f"vLLM 版本: {vllm.__version__}")  
  
# 测试GPU计算  
if torch.cuda.is_available():  
    device = "cuda"  
    x = torch.tensor([1.0, 2.0, 3.0]).to(device)  
    y = x * 2  
    print(f"GPU计算测试成功: {y}")  
else:  
    print("CUDA不可用!")
```

3. CUDA编程基础

CUDA相关的编程编译流程可以总结为如下流程图：



"Hello World"样例代码:

```

// add.cu
#include <iostream>
#include <math.h>

// Kernel function to add the elements of two arrays
__global__
void add(int n, float *x, float *y)
{
    for (int i = 0; i < n; i++)
        y[i] = x[i] + y[i];
}

int main(void)
{
    int N = 1<<20;
    float *x, *y;

    // Allocate Unified Memory – accessible from CPU or GPU
    cudaMallocManaged(&x, N*sizeof(float));
    cudaMallocManaged(&y, N*sizeof(float));

    // initialize x and y arrays on the host
    for (int i = 0; i < N; i++) {
        x[i] = 1.0f;
        y[i] = 2.0f;
    }
}
  
```

```

// Run kernel on 1M elements on the GPU
add<<<1, 1>>>(N, x, y);

// Wait for GPU to finish before accessing on host
cudaDeviceSynchronize();

// Check for errors (all values should be 3.0f)
float maxError = 0.0f;
for (int i = 0; i < N; i++) {
    maxError = fmax(maxError, fabs(y[i]-3.0f));
}
std::cout << "Max error: " << maxError << std::endl;

// Free memory
cudaFree(x);
cudaFree(y);
return 0;
}

```

(1) Kernel Function: 设备侧函数主要以__global__关键词修饰

```

__global__ add(int n, float* x, float* y) {
    for (int i = 0; i < n; ++i) {
        y[i] = x[i] + y[i];
    }
}

```

(2) Unified Memory: host-device统一编址的显存分配，利用cudaMallocManaged()分配统一内存，返回可访问的指针，kernel执行完毕需要手动利用cudaFree()释放内存（cudaMallocManaged-cudaFree和标准C++中的new-delete对应）

```

// Allocate Unified Memory -- accessible from CPU or GPU
float *x, *y, *sum;
cudaMallocManaged(&x, N*sizeof(float));
cudaMallocManaged(&y, N*sizeof(float));

...

// Free memory
cudaFree(x);
cudaFree(y);

```

(3) Kernal Launch: 从host侧启动Kernel函数，使用cuda的三重角度括号语法 <<<Dg, Db, Ns, S>>>

注释:

(a) Dg 代表整个grid的尺寸（一个grid有多少个block），数据类型为dim3，例如： $Dim3\ Dg(Dg.x, Dg.y, 1)$ 代表grid中每行有 $Dg.x$ 个block，每列有 $Dg.y$ 个block，第三个纬度恒定为1，因此整个grid一共有 $Dg.x * Dg.y$ 个block，且 $Dg.x, Dg.y \leq 65535$;

(b) Db 定义一个block的尺寸（一个block有多少thread），数据类型为dim3，例如： $Dim3\ Db(Db.x, Db.y, Db.z)$ 代表block中每行有 $Db.x$ 个thread，每列有 $Db.y$ 个thread，高度方向有 $Db.z$ 个thread，因此一个block有 $Db.x * Db.y * Db.z$ 个thread.

(c) Ns 为可选参数，用于设置每个block除了静态分配的共享内存外，最多能动态分配的共享内存大小，单位为Byte。如果不需要，则 $Ns = 0$ 或者参数缺省.

(d) S 为 $cudaStream_t$ 类型的可选参数，默认值为 $S = 0$ ，标识核函数位于哪个stream中（指定核函数在哪个stream中执行）

```
// 代表一个block，一个thread
add<<<1, 1>>>(N, sum, x, y);
```

(4) nvcc编译

- 基础编译命令：直接得到可执行文件

```
nvcc -o add add.cu
```

- 指定平台：从V100, A100, H100到B200不同GPU型号，nvcc可以编译得到特定平台二进制、ptx指令或者通用平台指令（就是各个平台的指令打包）

```
# 查询当前使用GPU平台的卡型号等信息
nvidia-smi
# 查询当前使用的GPU平台的计算能力
nvidia-smi --query-gpu=compute_cap --format=csv
```

编译到特定平台

(1) V100及所有计算能力7.0的GPU

```
nvcc -arch=sm_70 -o your_program your_program.cu
```

如果需要包含PTX代码以支持未来兼容性

```
nvcc -arch=sm_70 -gencode arch=compute_70,code=sm_70 -o your_program
your_program.cu
```

(2) A100及所有计算能力8.0的GPU

```
nvcc -arch=sm_80 -o your_program your_program.cu
```

如果需要包含PTX代码以支持未来兼容性

```
nvcc -arch=sm_80 -gencode arch=compute_80,code=sm_80 -o your_program  
your_program.cu
```

(3) H100及所有计算能力9.0的GPU

```
nvcc -arch=sm_90 -o your_program your_program.cu
```

如果需要包含PTX代码以支持未来兼容性

```
nvcc -arch=sm_90 -gencode arch=compute_90,code=sm_90 -o your_program  
your_program.cu
```

(4) B200及所有计算能力9.0的GPU

```
nvcc -arch=sm_90 -o your_program your_program.cu
```

如果需要包含PTX代码以支持未来兼容性

```
nvcc -arch=sm_90 -gencode arch=compute_90,code=sm_90 -o your_program  
your_program.cu
```

编译到通用平台

支持V100到H100的通用二进制

```
nvcc -gencode arch=compute_70,code=sm_70 -gencode  
arch=compute_80,code=sm_80 -gencode arch=compute_90,code=sm_90 -o  
universal_program program.cu
```

当前使用的GPU硬件为A100，因此以上add.cu代表编译命令为：

```
nvcc -arch=sm_80 -o add add.cu
```

执行后得到结果：

```
(pytorch) lthpc@gnode02:~/chengl/Programming> ./add  
Max error: 0
```

(5) cuda原生性能Profiling采集工具: nsys

- 直接采集完整性能统计数据

```
nsys profile -t cuda --stats=true ./add
```

结果会生成如下详细性能信息：包括内存分配（cudaMallocManaged），同步，内存释放，Kernel启动时间，Kernel执行时间等

```
(pytorch) lthpc@gnode02:~/chengl/Programming> nsys profile -t cuda --  
stats=true ./add  
Max error: 0
```

Generating '/tmp/nsys-report-94b0.qdstrm'

[1/6] [=====100%] report2.nsys-rep

[2/6] [=====100%] report2.sqlite

[3/6] Executing 'cuda_api_sum' stats report

Time (%)	Total Time (ns)	Num Calls	Avg (ns)	Med (ns)	
Min (ns)	Max (ns)	StdDev (ns)	Name		
64.2	308,780,682	2	154,390,341.0	154,390,341.0	
19,020	308,761,662	218,314,015.8	cudaMallocManaged		
35.7	171,461,980	1	171,461,980.0	171,461,980.0	
171,461,980	171,461,980	0.0	cudaDeviceSynchronize		
0.1	468,500	2	234,250.0	234,250.0	
183,820	284,680	71,318.8	cudaFree		
0.0	197,110	1	197,110.0	197,110.0	
197,110	197,110	0.0	cudaLaunchKernel		
0.0	840	1	840.0	840.0	
840	840	0.0	cuModuleGetLoadingMode		

[4/6] Executing 'cuda_gpu_kern_sum' stats report

Time (%)	Total Time (ns)	Instances	Avg (ns)	Med (ns)	
Min (ns)	Max (ns)	StdDev (ns)	Name		
100.0	171,458,500	1	171,458,500.0	171,458,500.0	
171,458,500	171,458,500	0.0	add(int, float *, float *)		

[5/6] Executing 'cuda_gpu_mem_time_sum' stats report

Time (%)	Total Time (ns)	Count	Avg (ns)	Med (ns)	Min (ns)	Max
(ns)	StdDev (ns)	Operation				
61.3	978,906	48	20,393.9	6,575.5	3,679	
105,411	29,701.5	[CUDA memcpy Unified Host-to-Device]				
38.7	617,794	24	25,741.4	5,167.5	2,655	
151,107	43,848.0	[CUDA memcpy Unified Device-to-Host]				

[6/6] Executing 'cuda_gpu_mem_size_sum' stats report

Total (MB)	Count	Avg (MB)	Med (MB)	Min (MB)	Max (MB)	StdDev (MB)
Operation						

8.389	48	0.175	0.033	0.004	1.044	0.304
[CUDA memcpy Unified Host-to-Device]						
4.194	24	0.175	0.033	0.004	1.044	0.307
[CUDA memcpy Unified Device-to-Host]						

Generated:

/home/lthpc/chengl/Programming/report2.nsys-rep
/home/lthpc/chengl/Programming/report2.sqlite

- 简化性能统计信息：Mark Harris 写了一个简单的nsys_easy脚本(https://github.com/harrism/nsys_easy)，用来简化输出信息

```
# 修改执行权限
chmod 755 ~/nsys_easy/nsys_easy
# 添加nsys_easy环境变量
export PATH=~/nsys_easy:$PATH
# 采集性能信息
nsys_easy ./add
```

结果会生成如下简化信息：

```
(pytorch) lthpc@gnode02:~/chengl/Programming> nsys_easy ./add
Max error: 0
Generating '/tmp/nsys-report-66ad.qdstrm'
[1/1] [=====100%] nsys_easy.nsys-rep
Generated:
    /home/lthpc/chengl/Programming/nsys_easy.nsys-rep
Generating SQLite file nsys_easy.sqlite from nsys_easy.nsys-rep
Processing 939 events:
[=====100%]
Processing [nsys_easy.sqlite] with [/home/software/cuda-12.3/nsight-
systems-2023.3.3/host-linux-x64/reports/cuda_gpu_sum.py]...
```

** CUDA GPU Summary (Kernels/MemOps) (cuda_gpu_sum):

Time (%)	Total Time (ns)	Instances	Avg (ns)	Med (ns)
Min (ns)	Max (ns)	StdDev (ns)	Category	
Operation				

-----	-----	-----	-----	-----
-----	-----	-----	-----	-----

```

    99.1      172,191,985      1  172,191,985.0  172,191,985.0
172,191,985  172,191,985      0.0  CUDA_KERNEL  add(int, float *,
float *)
    0.6      979,518      48      20,406.6      6,527.5
3,647      105,666      29,680.1  MEMORY_OPER  [CUDA memcpy Unified
Host-to-Device]
    0.4      618,691      24      25,778.8      5,199.5
2,495      151,235      43,882.6  MEMORY_OPER  [CUDA memcpy Unified
Device-to-Host]

```

(6) PyTorch + CUDA原生集成开发 + pyTorch统一性能Profiling采集工具： torch.profiler

参考PyTorch官网说明：<https://docs.pytorch.org/docs/stable/profiler.html>

目前LLM基本都是在PyTorch的pythonic环境下开发使用，为了实现CUDA函数在PyTorch的自然集成，需要用到`torch.utils.cpp_extension`和`torch/extension.h`，并利用`pybind11`完成pytorch下对CUDA-C/C++函数的调用。<https://github.com/gau-nernst/learn-cuda>中给了10个示例教程，这里给出01-vector addition的代码示例。

Name	Description
01. Vector addition	Simple example to get everything working.
02a. Matrix multiplication SIMT	Block tiling, thread tiling, warp tiling.
02b. Matrix multiplication TensorOp	Inline PTX, <code>cvta</code> , <code>ldmatrix</code> , <code>mma</code> .
03. Sum	Reduction in general. Prepare for softmax (max and sum).
04. Softmax	Naive (safe) softmax, online softmax. <code>atomicCAS()</code> . Single-block and multi-block per row.
05. FP6	FP6 primitives (FP32/FP16/BF16<->FP6).
06. Box blur	2D CUDA blocks/threads. TODO: optimize with separable filters, moving average.
07. Attention	Flash attention
08. Row-scaled matmul	Simple epilogue
09. Block-scaled matmul	MXFP8

```

// add.cu
#include <torch/extension.h>

```

```

#define CHECK_CUDA(x) TORCH_CHECK(x.device().is_cuda(), #x " must be a
CUDA tensor")
#define CHECK_CONTIGUOUS(x) TORCH_CHECK(x.is_contiguous(), #x " must be
contiguous")
#define CHECK_INPUT(x)

CHECK_CUDA(x);

CHECK_CONTIGUOUS(x)

__global__ void add_kernel(const float *input1, const float *input2,
float *output, int size) {
    const int idx = blockIdx.x * blockDim.x + threadIdx.x;
    if (idx < size)
        output[idx] = input1[idx] + input2[idx];
}

torch::Tensor add(torch::Tensor input1, torch::Tensor input2) {
    CHECK_INPUT(input1);
    CHECK_INPUT(input2);
    int size = input1.numel();
    TORCH_CHECK(size == input2.numel(), "input1 and input2 must have the
same size");
    torch::Tensor output = torch::empty(size, input1.options());

    int n_threads = 256;
    int n_blocks = (size + n_threads - 1) / n_threads;
    add_kernel<<<n_blocks, n_threads>>>(input1.data_ptr<float>(),
input2.data_ptr<float>(), output.data_ptr<float>(), size);

    return output;
}

PYBIND11_MODULE(TORCH_EXTENSION_NAME, m) { m.def("add", &add, "Add two
vectors"); }

```

实际上就是添加了一个tensor的wrapper，因为tensor是pytorch的基础数据类型，因此核心步骤就是将原始cuda的入参和返回参数等转成tensor（指针），传给torch实现包装。下面是pytorch中的调用代码（可以看到就是重新声明了cuda函数，将其归类为一类module，从而下面可以直接从module中调用原始封装后的cuda函数，如果需要将函数注册进torch，即torch.add，后续会涉及）：

```

# main.py
import torch

```

```

import torch.utils.cpp_extension
from torch.profiler import profile, record_function, ProfilerActivity
import matplotlib.pyplot as plt

module = torch.utils.cpp_extension.load(
    "module",
    sources=["add.cu"],
    extra_cuda_cflags=["-O3"],
    verbose=True,
)

# Example usage
input1 = torch.randn(1024000, device="cuda")
input2 = torch.randn(1024000, device="cuda")
output = module.add(input1, input2)

```

下面给出torch.profiler采集以上封装后的cuda函数方法。

- (a) 最基础的torch.profiler裸测试调用采集性能数据

```

# 核心结构
from torch.profiler import profile, record_function, ProfilerActivity
with torch.profiler.profile(
    activities=[
        torch.profiler.ProfilerActivity.CPU,
        torch.profiler.ProfilerActivity.CUDA,
    ]
) as p:
    code_to_profile() # 替换成需要采集的torch代码 (函数)
print(p.key_averages().table(sort_by="self_cuda_time_total",
row_limit=-1))

```

完整测试代码如下：

```

import torch
import torch.utils.cpp_extension
from torch.profiler import profile, record_function, ProfilerActivity
import matplotlib.pyplot as plt

module = torch.utils.cpp_extension.load(
    "module",
    sources=["add.cu"],
    extra_cuda_cflags=["-O3"],
    verbose=True,
)

```

```

)

# Example usage
input1 = torch.randn(1024000, device="cuda")
input2 = torch.randn(1024000, device="cuda")

with torch.profiler.profile(
    activities=[
        torch.profiler.ProfilerActivity.CPU,
        torch.profiler.ProfilerActivity.CUDA,
    ]
) as p:
    output = module.add(input1, input2)
print(p.key_averages().table(sort_by="self_cuda_time_total",
row_limit=-1))
p.export_chrome_trace("trace.json") # 可以生成程序运行的timeline

```

运行后的结果如下：

```

(pytorch) lthpc@gnode02:~/chengl/Programming/learn-
cuda/1_vector_addition> CUDA_VISIBLE_DEVICES=0 python main.py
Using /home/lthpc/.cache/torch_extensions/py310_cu126 as PyTorch
extensions root...
Detected CUDA files, patching ldflags
Emitting ninja build file
/home/lthpc/.cache/torch_extensions/py310_cu126/module/build.ninja...
Building extension module module...
Allowing ninja to set a default number of workers... (overridable by
setting the environment variable MAX_JOBS=N)
ninja: no work to do.
Loading extension module module...

```

Self CPU %	CPU total %	CPU total	CPU time avg	Name	Self CPU %
Self CUDA %	CUDA total	CUDA time avg	# of Calls		Self CUDA
100.00%	9.824us	9.824us	1	add_kernel(float const*, float const*, float*, int)	0.00%
					9.824us

			aten::empty	5.96%
52.520us	75.63%	666.350us	666.350us	0.000us
0.00%	0.000us	0.000us	1	
			Unrecognized	69.67%
613.830us	69.67%	613.830us	613.830us	0.000us
0.00%	0.000us	0.000us	1	
			cudaLaunchKernel	8.46%
74.530us	23.81%	209.810us	209.810us	0.000us
0.00%	0.000us	0.000us	1	
			Runtime Triggered Module Loading	15.35%
135.280us	15.35%	135.280us	67.640us	0.000us
0.00%	0.000us	0.000us	2	
			cudaDeviceSynchronize	0.56%
4.910us	0.56%	4.910us	4.910us	0.000us
0.00%	0.000us	0.000us	1	

Self CPU time total: 881.070us
Self CUDA time total: 9.824us

从上面列表中找到Name = add_kernel对应的Self CUDA= 9.824us即为GPU上add_kernel执行的完整时间，下面还可以发现cudaLaunchKernel代表从host侧启动GPU侧核函数花费时间为209.810us，可以看到远远长于Kernel本身在GPU上执行时间。

- (b) 考虑预热等因素后的torch.profiler调用性能采集方法

```
# 核心结构
with torch.profiler.profile(
    activities=[
        torch.profiler.ProfilerActivity.CPU,
        torch.profiler.ProfilerActivity.CUDA,
    ]
) as p:
    code_to_profile_0()
    // turn off collection of all CUDA activity
    p.toggle_collection_dynamic(False,
[torch.profiler.ProfilerActivity.CUDA])
    code_to_profile_1()
    // turn on collection of all CUDA activity
    p.toggle_collection_dynamic(True,
[torch.profiler.ProfilerActivity.CUDA])
    code_to_profile_2()
```

```
print(p.key_averages().table(sort_by="self_cuda_time_total",
row_limit=-1))
```

完整测试代码:

```
import torch
import torch.utils.cpp_extension
from torch.profiler import profile, record_function, ProfilerActivity
import matplotlib.pyplot as plt

module = torch.utils.cpp_extension.load(
    "module",
    sources=["add.cu"],
    extra_cuda_cflags=["-O3"],
    verbose=True,
)

# Example usage
input1 = torch.randn(1024000, device="cuda")
input2 = torch.randn(1024000, device="cuda")

# Non-default profiler schedule allows user to turn profiler on and off
# on different iterations of the training loop;
# trace_handler is called every time a new trace becomes available
def trace_handler(prof):
    print(
        prof.key_averages().table(sort_by="self_cuda_time_total",
row_limit=-1)
    )
    prof.export_chrome_trace("/tmp/test_trace_" + str(prof.step_num) +
".json")

with torch.profiler.profile(
    activities=[
        torch.profiler.ProfilerActivity.CPU,
        torch.profiler.ProfilerActivity.CUDA,
    ],
    # In this example with wait=1, warmup=1, active=2, repeat=1,
    # profiler will skip the first step/iteration,
    # start warming up on the second, record
    # the third and the forth iterations,
    # after which the trace will become available
    # and on_trace_ready (when set) is called;
    # the cycle repeats starting with the next step
```

```

    schedule=torch.profiler.schedule(wait=1, warmup=1, active=7,
repeat=1),
    on_trace_ready=trace_handler,
    record_shapes=True,
    # on_trace_ready=torch.profiler.tensorboard_trace_handler('./log')
    # used when outputting for tensorboard
) as p:
    for iter in range(10):
        output = module.add(input1, input2)
        # send a signal to the profiler that the next iteration has
started
        p.step()

```

采集结果如下：

```

(pytorch) lthpc@gnode02:~/chengl/Programming/learn-
cuda/1_vector_addition> CUDA_VISIBLE_DEVICES=0 python main.py
Using /home/lthpc/.cache/torch_extensions/py310_cu126 as PyTorch
extensions root...
Detected CUDA files, patching ldflags
Emitting ninja build file
/home/lthpc/.cache/torch_extensions/py310_cu126/module/build.ninja...
Building extension module module...
Allowing ninja to set a default number of workers... (overridable by
setting the environment variable MAX_JOBS=N)
ninja: no work to do.
Loading extension module module...

```

Self CPU Self CUDA %	CPU total % CUDA total	CPU total CUDA time avg	CPU time avg # of Calls	Name Self CPU % Self CUDA
0.000us	0.00%	0.000us	0.000us	add_kernel(float const*, float const*, float*, int)
100.00%	45.956us	6.565us	7	
0.000us	0.00%	0.000us	0.000us	ProfilerStep*
100.00%	45.956us	6.565us	7	

			ProfilerStep*	71.61%
380.100us	96.98%	514.710us	73.530us	0.000us
0.00%	0.000us	0.000us	7	
			aten::empty	12.90%
68.450us	12.90%	68.450us	9.779us	0.000us
0.00%	0.000us	0.000us	7	
			cudaLaunchKernel	12.47%
66.160us	12.47%	66.160us	9.451us	0.000us
0.00%	0.000us	0.000us	7	
			cudaDeviceSynchronize	3.02%
16.050us	3.02%	16.050us	16.050us	0.000us
0.00%	0.000us	0.000us	1	

Self CPU time total: 530.760us

Self CUDA time total: 45.956us

可以看到考虑预热和多次平均后的add_kernel时间为6.565us，并且多次运行时间差别也不大。

(c) 考虑预热等因素 + 仅采集部分代码的torch.profiler调用性能采集方法：
toggle_collection_dynamic

```
with torch.profiler.profile(
    activities=[
        torch.profiler.ProfilerActivity.CPU,
        torch.profiler.ProfilerActivity.CUDA,
    ]
) as p:
    code_to_profile_0()
    // turn off collection of all CUDA activity
    p.toggle_collection_dynamic(False,
[torch.profiler.ProfilerActivity.CUDA])
    code_to_profile_1()
    // turn on collection of all CUDA activity
    p.toggle_collection_dynamic(True,
[torch.profiler.ProfilerActivity.CUDA])
    code_to_profile_2()
print(p.key_averages().table(sort_by="self_cuda_time_total",
row_limit=-1))
```

完整代码如下：


```

import torch
import torch.utils.cpp_extension
from torch.profiler import profile, record_function, ProfilerActivity
import matplotlib.pyplot as plt

module = torch.utils.cpp_extension.load(
    "module",
    sources=["add.cu"],
    extra_cuda_cflags=["-O3"],
    verbose=True,
)

# Non-default profiler schedule allows user to turn profiler on and off
# on different iterations of the training loop;
# trace_handler is called every time a new trace becomes available
def trace_handler(prof):
    print(
        prof.key_averages().table(sort_by="self_cuda_time_total",
row_limit=-1)
    )
    prof.export_chrome_trace("/tmp/test_trace_" + str(prof.step_num) +
".json")

with torch.profiler.profile(
    activities=[
        torch.profiler.ProfilerActivity.CPU,
        torch.profiler.ProfilerActivity.CUDA,
    ],
    # In this example with wait=1, warmup=1, active=2, repeat=1,
    # profiler will skip the first step/iteration,
    # start warming up on the second, record
    # the third and the forth iterations,
    # after which the trace will become available
    # and on_trace_ready (when set) is called;
    # the cycle repeats starting with the next step
    schedule=torch.profiler.schedule(wait=1, warmup=1, active=7,
repeat=1),
    on_trace_ready=trace_handler,
    record_shapes=True,
    # on_trace_ready=torch.profiler.tensorboard_trace_handler('./log')
    # used when outputting for tensorboard
) as p:
    for iter in range(10):

```

```
p.toggle_collection_dynamic(False,
[torch.profiler.ProfilerActivity.CPU,
torch.profiler.ProfilerActivity.CUDA]) # 同时关掉CPU和GPU上的性能采集
input1 = torch.randn(1024000, device="cuda") # 为了测试效果，将数
据生成也考虑如下采集
input2 = torch.randn(1024000, device="cuda")
p.toggle_collection_dynamic(True,
[torch.profiler.ProfilerActivity.CUDA]) # 仅仅打开GPU上的性能采集
output = module.add(input1, input2)
# send a signal to the p
p.step()
```

profiling结果如下:

[illegible]

作为对比：

(a) 类型一

```
for iter in range(10):
    p.toggle_collection_dynamic(False,
    [torch.profiler.ProfilerActivity.CPU,
    torch.profiler.ProfilerActivity.CUDA]) # 同时关掉CPU和GPU上的性能采集
    input1 = torch.randn(1024000, device="cuda") # 为了测试效果，将数
    据生成也考虑如下采集
    input2 = torch.randn(1024000, device="cuda")
    output = module.add(input1, input2)
    # send a signal to the p
    p.step()
```

```

Ninja: no work to do.
Loading extension module module...
-----
      Name      Self CPU %      Self CPU      CPU total %      CPU total      CPU time avg      # of Calls
-----
  ProfilerStep*  100.00%    127.950us    100.00%    127.950us    127.950us          1
-----
Self CPU time total: 127.950us

```

(b) 类型二

```

for iter in range(10):
    p.toggle_collection_dynamic(False,
    [torch.profiler.ProfilerActivity.CPU,
    torch.profiler.ProfilerActivity.CUDA]) # 同时关掉CPU和GPU上的性能采集
    input1 = torch.randn(1024000, device="cuda") # 为了测试效果，将数
据生成也考虑如下采集
    input2 = torch.randn(1024000, device="cuda")
    p.toggle_collection_dynamic(True,
    [torch.profiler.ProfilerActivity.CPU]) # 仅仅打开CPU上的性能采集
    output = module.add(input1, input2)
    # send a signal to the p
    p.step()

```

Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	# of Calls
ProfilerStep*	96.15%	900.710us	100.00%	936.750us	133.821us	7
aten::empty	3.85%	36.040us	3.85%	36.040us	5.149us	7

Self CPU time total: 936.750us

(c) 类型三

```

for iter in range(10):
    p.toggle_collection_dynamic(False,
    [torch.profiler.ProfilerActivity.CPU,
    torch.profiler.ProfilerActivity.CUDA]) # 同时关掉CPU和GPU上的性能采集
    input1 = torch.randn(1024000, device="cuda") # 为了测试效果，将数
据生成也考虑如下采集
    input2 = torch.randn(1024000, device="cuda")
    p.toggle_collection_dynamic(True,
    [torch.profiler.ProfilerActivity.CPU,
    torch.profiler.ProfilerActivity.CUDA]) # 打开CPU+GPU上的性能采集
    output = module.add(input1, input2)
    # send a signal to the p
    p.step()

```

Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	Self CUDA	Self CUDA %	CUDA total	CUDA time avg	# of Calls
add_kernel(float const*, float const*, float*, int)	0.00%	0.000us	0.00%	0.000us	0.000us	41.474us	100.00%	41.474us	5.925us	7
ProfilerStep*	0.00%	0.000us	0.00%	0.000us	0.000us	41.474us	100.00%	41.474us	5.925us	7
ProfilerStep*	91.86%	17.944ms	92.29%	18.029ms	2.576ms	0.000us	0.00%	0.000us	0.000us	7
aten::empty	0.19%	37.720us	0.19%	37.720us	5.389us	0.000us	0.00%	0.000us	0.000us	7
cudaLaunchKernel	0.24%	47.080us	0.24%	47.080us	6.726us	0.000us	0.00%	0.000us	0.000us	7
cudaDeviceSynchronize	7.71%	1.506ms	7.71%	1.506ms	1.506ms	0.000us	0.00%	0.000us	0.000us	1

Self CPU time total: 19.535ms
Self CUDA time total: 41.474us

(d) 类型四

```

for iter in range(10):
    p.toggle_collection_dynamic(True,
    [torch.profiler.ProfilerActivity.CPU,
    torch.profiler.ProfilerActivity.CUDA]) # 同时关掉CPU和GPU上的性能采集
    input1 = torch.randn(1024000, device="cuda") # 为了测试效果，将数据生成也考虑如下采集
    input2 = torch.randn(1024000, device="cuda")
    p.toggle_collection_dynamic(True,
    [torch.profiler.ProfilerActivity.CPU,
    torch.profiler.ProfilerActivity.CUDA]) # 打开CPU+GPU上的性能采集
    output = module.add(input1, input2)
    # send a signal to the p
    p.step()

```

	Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	Self CUDA	Self CUDA %	CUDA total	CUDA time avg	# of Calls
	ProfilerStep*	0.00%	0.000us	0.00%	0.000us	0.000us	374.888us	188.10%	374.888us	53.555us	7
	aten::normal_	10.47%	106.230us	23.78%	241.180us	17.227us	159.298us	79.93%	159.298us	11.378us	14
void at::native::(anonymous namespace)::distribution...		0.00%	0.000us	0.00%	0.000us	0.000us	159.298us	79.93%	159.298us	11.378us	14
add_kernel(float const*, float const*, float*, int)		0.00%	0.000us	0.00%	0.000us	0.000us	40.001us	20.07%	40.001us	5.714us	7
	ProfilerStep*	53.94%	547.360us	99.12%	1.005ms	143.636us	0.000us	0.00%	159.298us	22.757us	7
	aten::randn	7.24%	73.400us	38.33%	388.780us	27.770us	0.000us	0.00%	159.298us	11.378us	14
	aten::empty	10.20%	103.470us	10.20%	103.470us	4.927us	0.000us	0.00%	0.000us	0.000us	21
	cudaStreamIsCapturing	1.67%	16.980us	1.67%	16.980us	0.406us	0.000us	0.00%	0.000us	0.000us	28
	cudaLaunchKernel	15.58%	158.010us	15.58%	158.010us	7.524us	0.000us	0.00%	0.000us	0.000us	21
	cudaDeviceSynchronize	0.88%	8.880us	0.88%	8.880us	8.880us	0.000us	0.00%	0.000us	0.000us	1
Self CPU time total: 1.014ms											
Self CUDA time total: 199.299us											

(7) CUDA Kernel基本优化方法